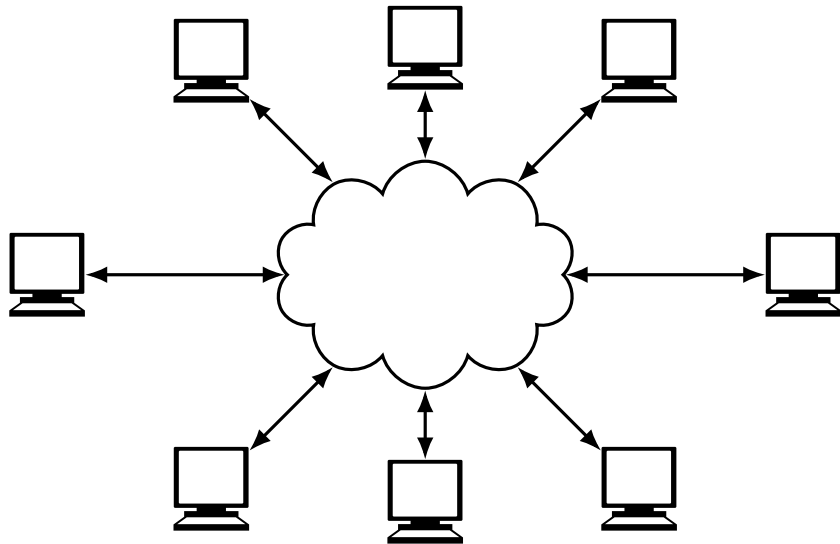
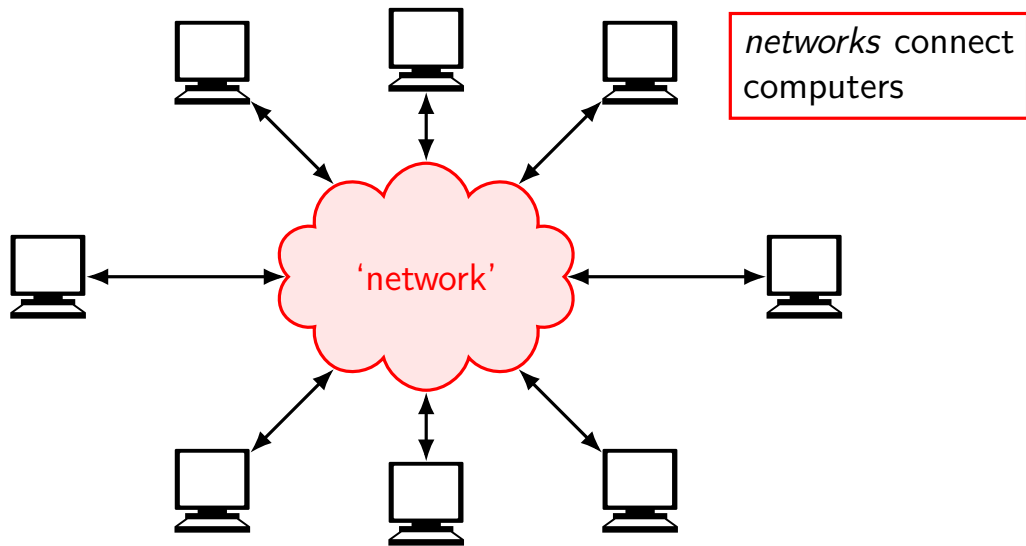




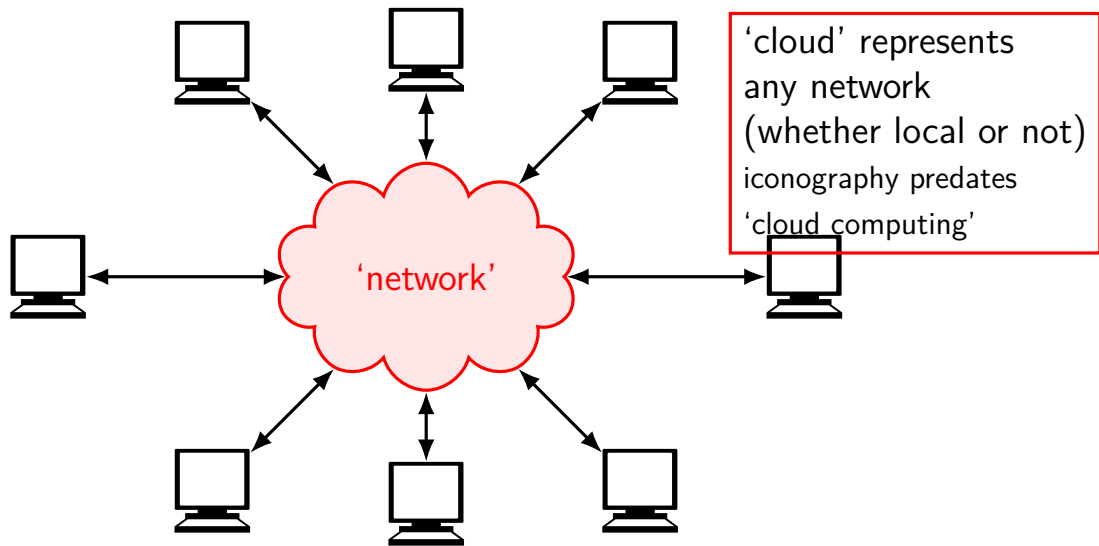
# networks / hosts aka end systems



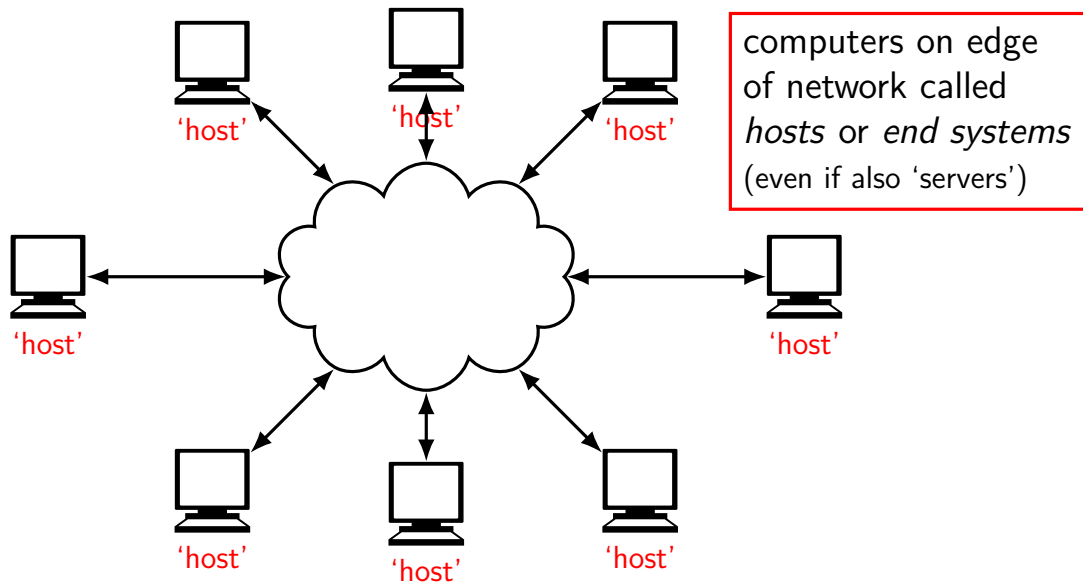
# networks / hosts aka end systems



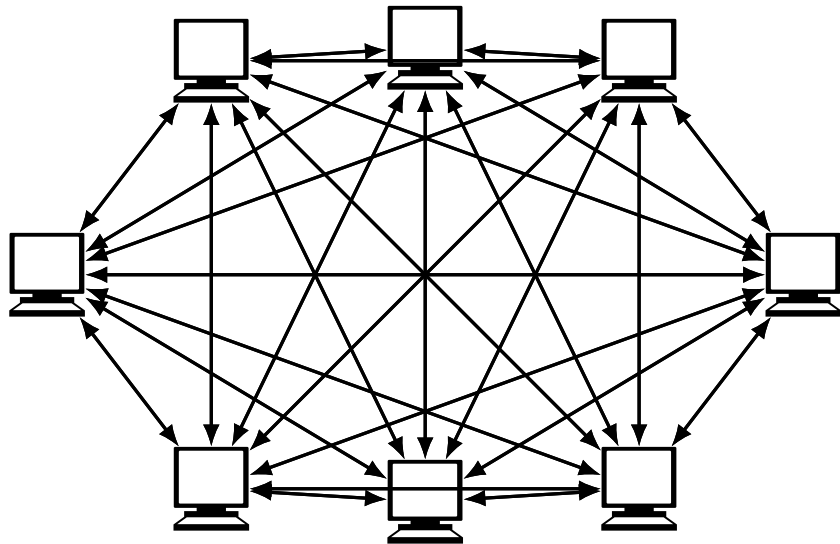
# networks / hosts aka end systems



# networks / hosts aka end systems



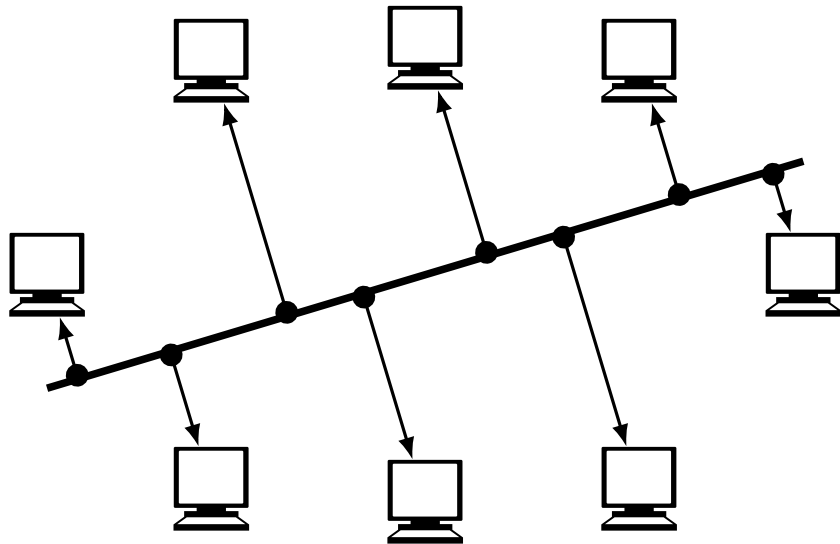
direct connections?



## shared medium: radio?

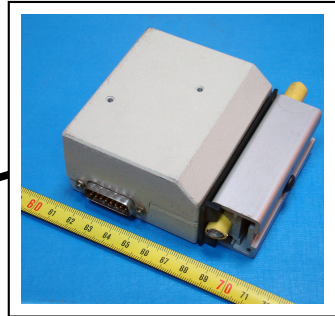
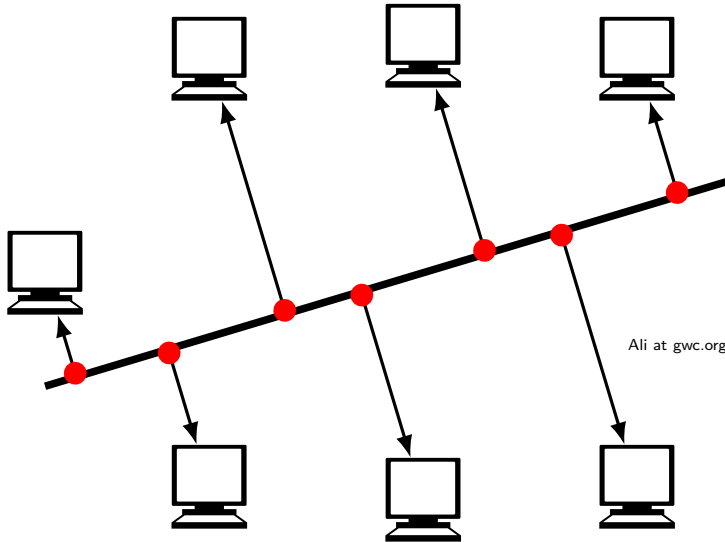


## shared medium: wires



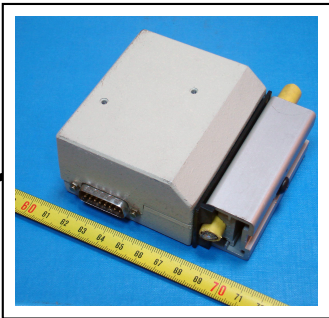



# shared medium: wires



Ali at gwc.org.uk / Alistair1978 via Wikimedia commons / CC-BY-SA 2.5

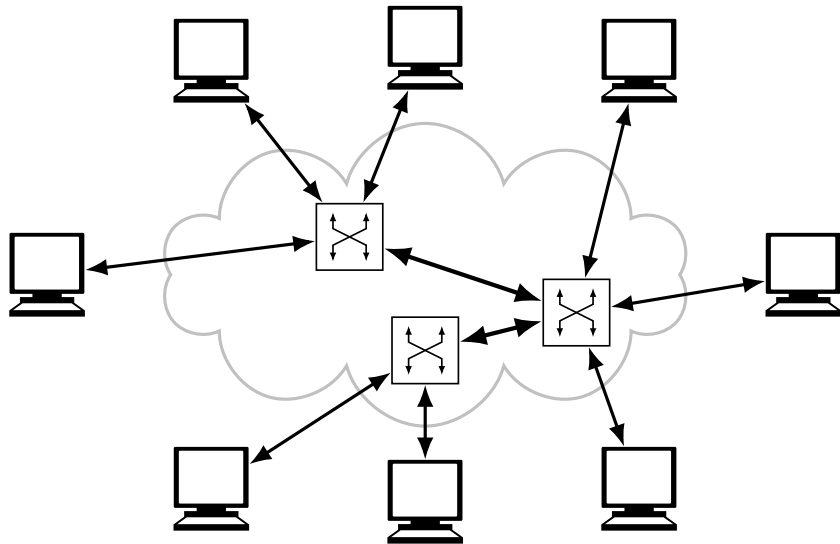
# shared medium: wires



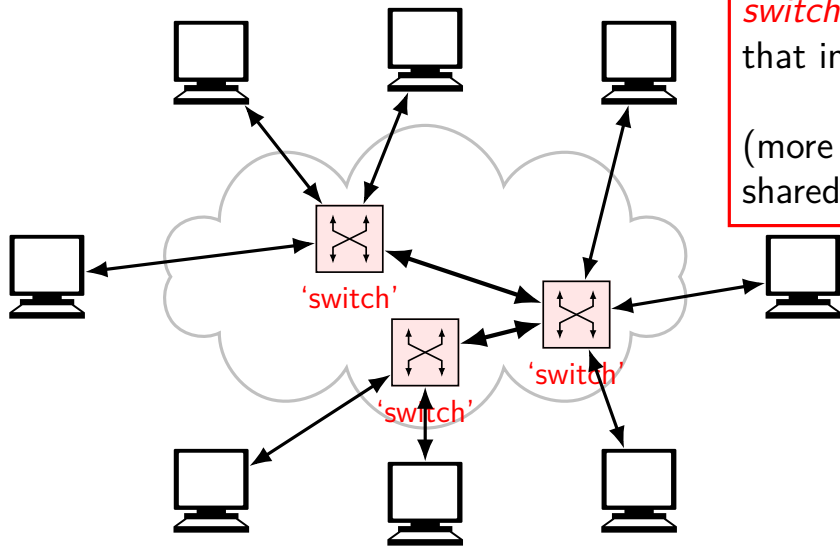
Ali at gwc.org.uk / Alistair1978 via Wikimedia commons / CC-BY-SA 2.5



# switches / nodes / links



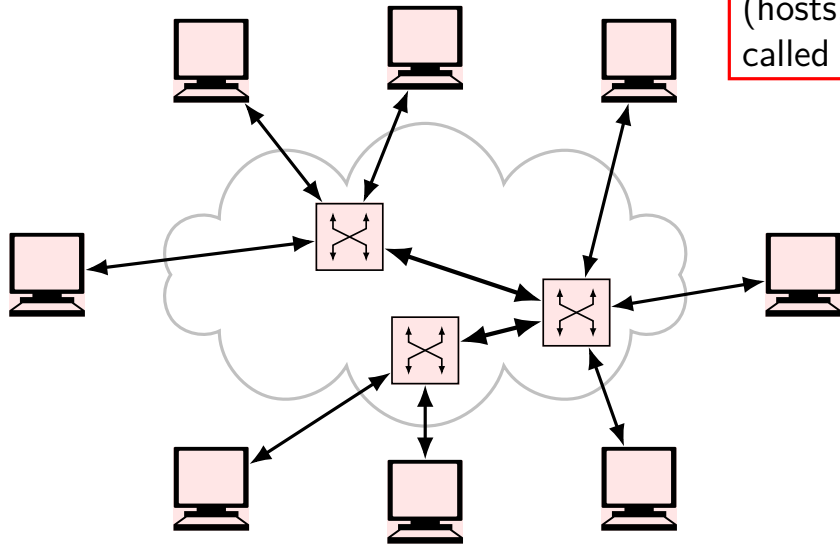
# switches / nodes / links



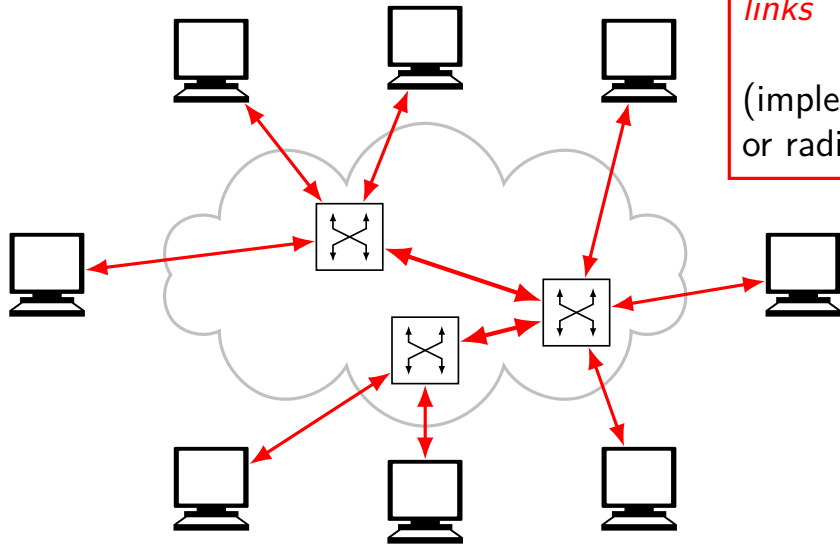
hosts directly connected to  
*switches*  
that implement network  
(more efficiently than  
shared medium)

# switches / nodes / links

machines on network  
(hosts, switches, ...)  
called *nodes*



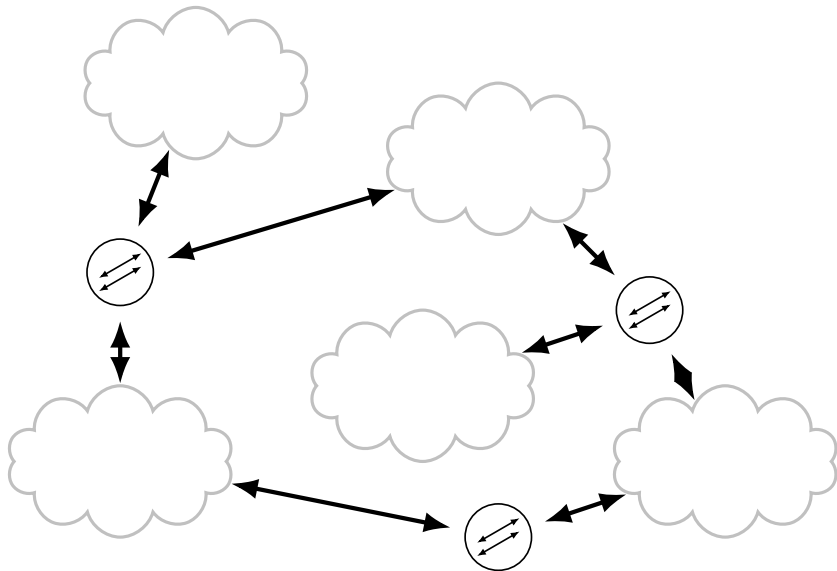
# switches / nodes / links



nodes connected by  
*links*

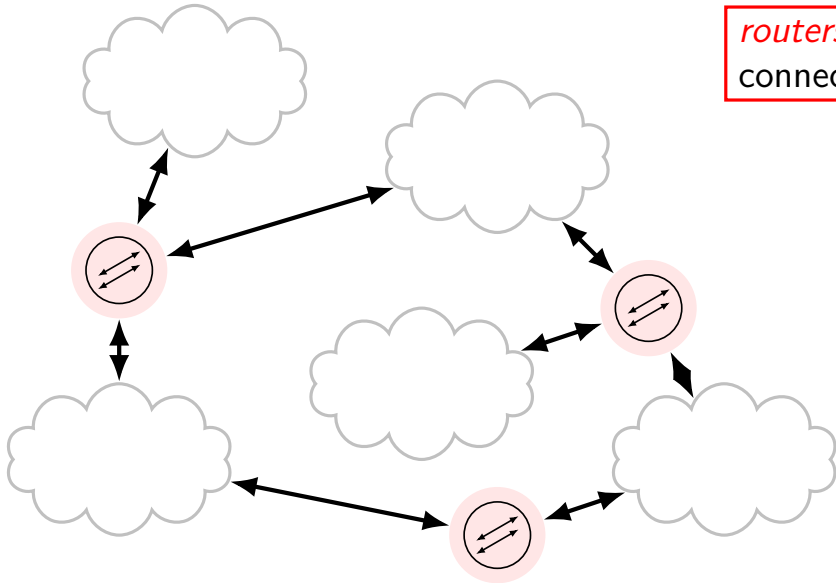
(implemented by wires  
or radio or ...)

# routers / internetwork



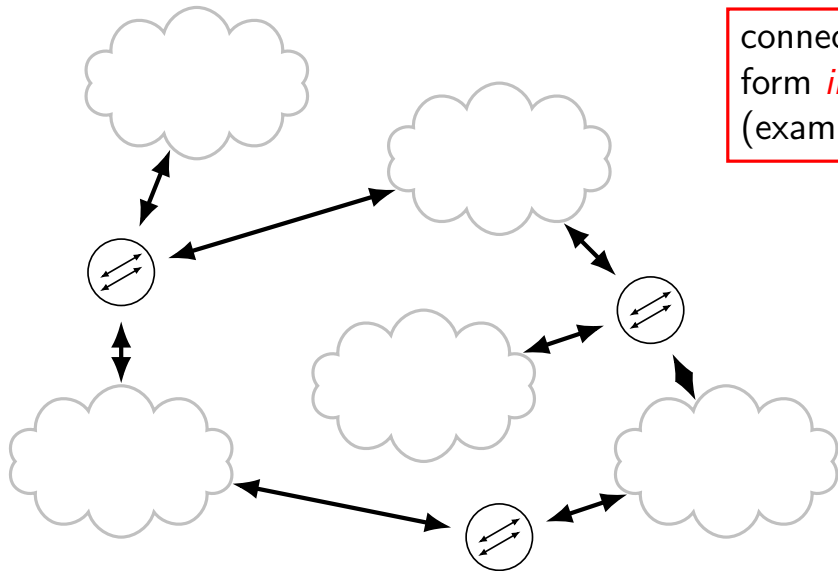
# routers / internetwork

*routers* or *gateways*  
connect networks



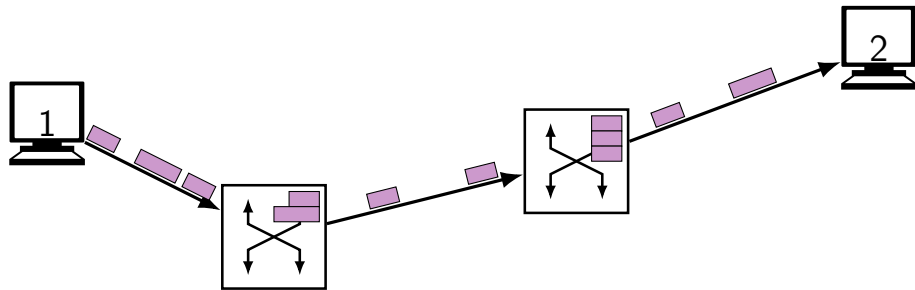


# routers / internetwork



connected networks  
form *internetwork*  
(example: the Internet)

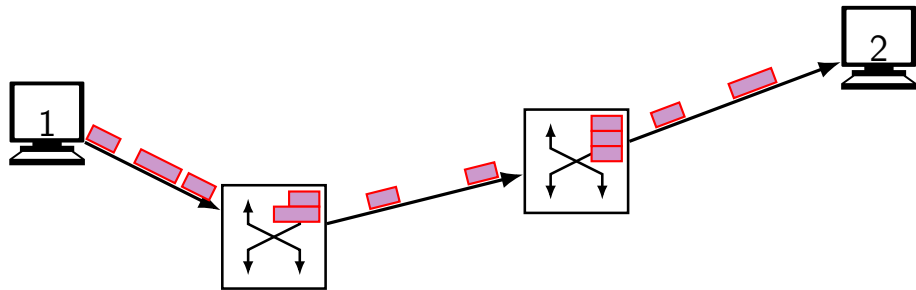
# flows / packets



*flow* of data between two machines

*flow* is very general term  
will depend on context how it relates to  
connections, sockets, etc.


# flows / packets



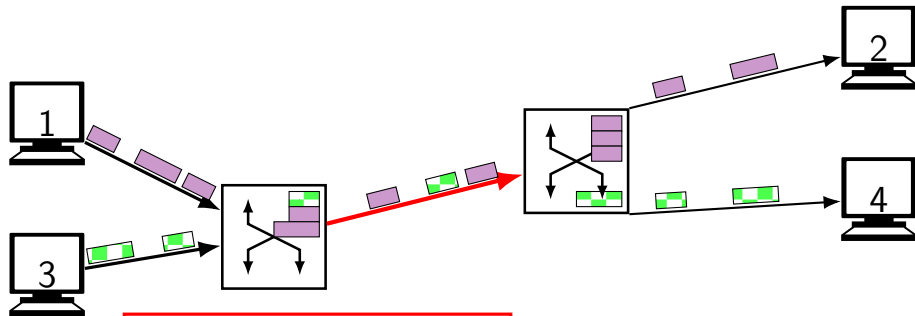
*flow* of data between two machines

possibly divided up into pieces,  
called *packets*, *frames*, *segments*  
(which name is best depends on context)

# (de)multiplexing

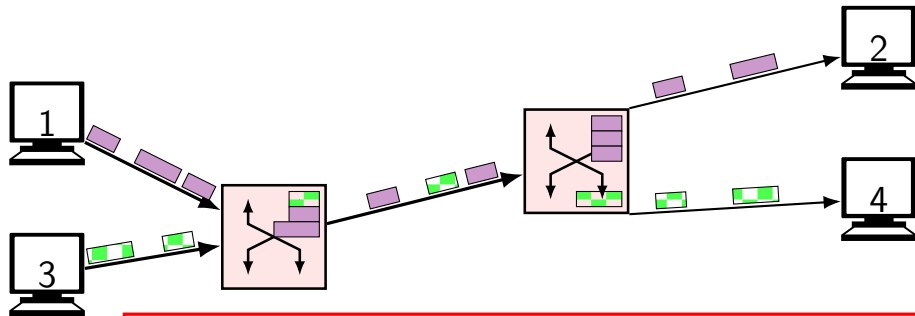


# (de)multiplexing



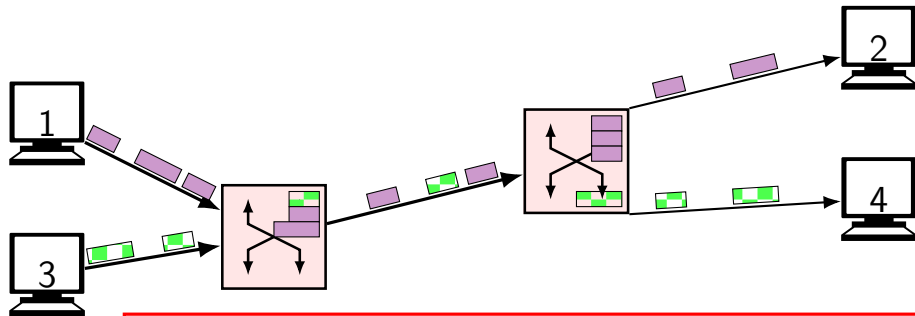
two or more flows can  
share one or more links

# (de)multiplexing



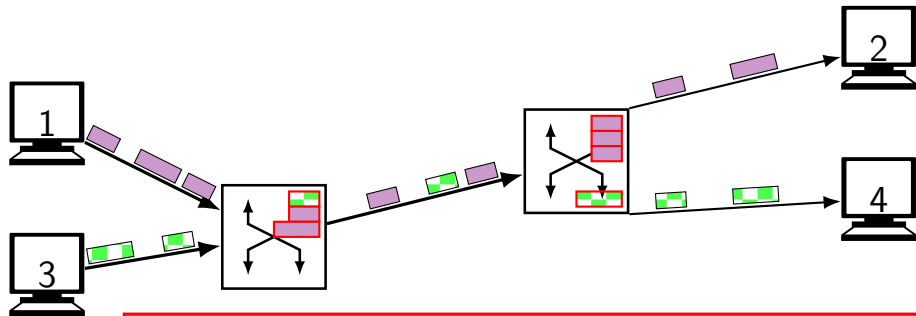
left switch *multiplexes* the two flows onto one link  
right switch *demultiplexes* them to separate them

# (de)multiplexing



this picture: multiplexed by dividing up *time* on link

# (de)multiplexing

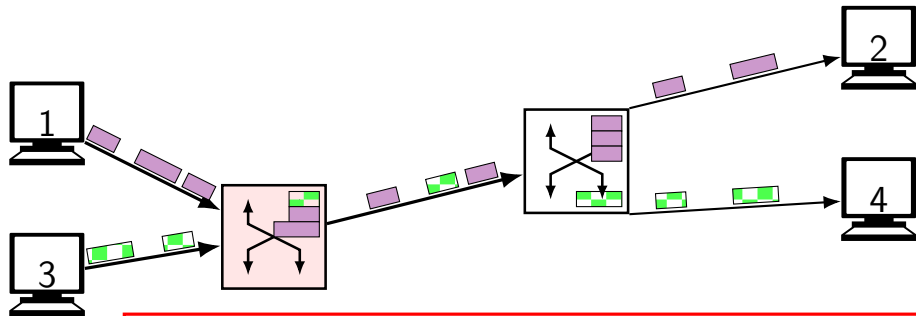


switches usually have *buffers* (also called *queues*)  
hold waiting packets

absorbs temporary “bursts” where packets come faster  
than outgoing link can handle



# (de)multiplexing




incomplete list of causes of 'bursts':

multiple unsynchronized flows

fast links produce packets faster for slow can send

# (de)multiplexing

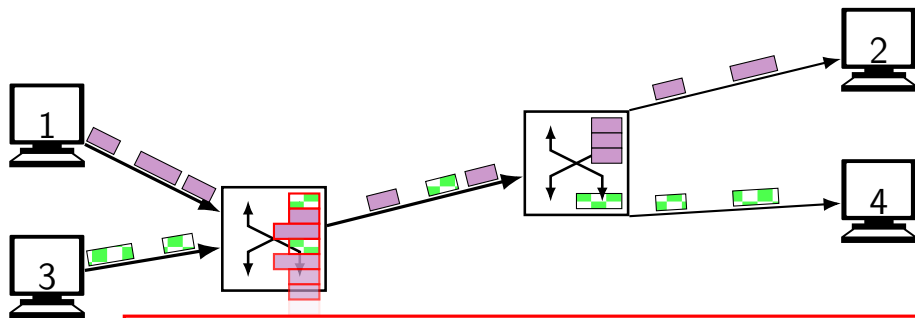


incomplete list of causes of 'bursts':

multiple unsynchronized flows


fast links produce packets faster for slow can send

# (de)multiplexing

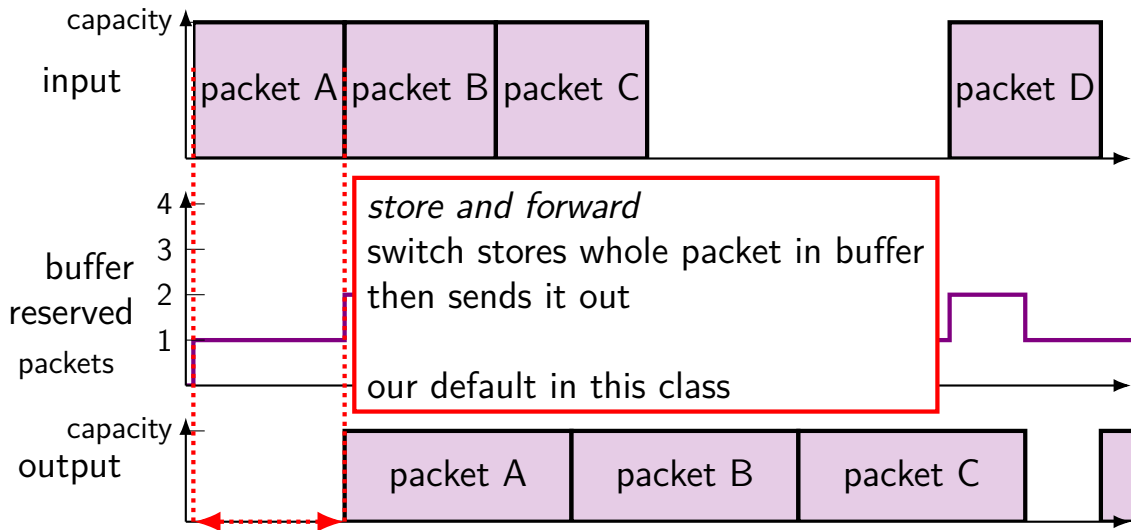


if buffer full, switch must *drop* packets  
will happen eventually if overall rate faster than outgoing link  
scenario is called *congestion*

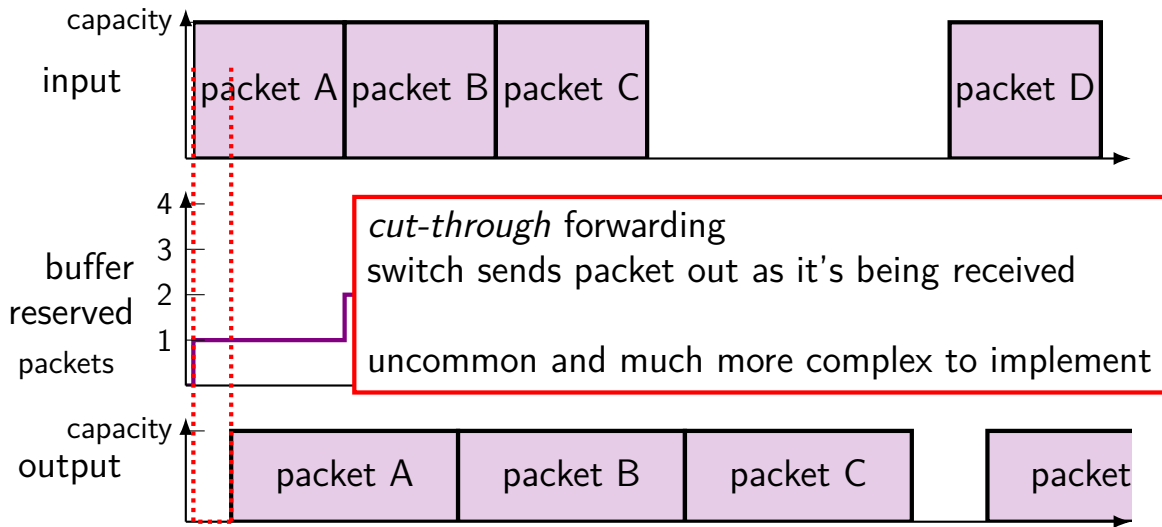
# buffer usage: fast to slow, store + forward



# buffer usage: fast to slow, store + forward



# buffer usage: fast to slow, cut-through



# channel abstractions

want to avoid custom network for each application

but applications have different needs

→ multiple application interfaces to networks

common implementation of **common patterns**

# some abstractions

## stream

continuous stream of bytes from one program to another  
'connection' from one program to another

## datagrams

send small messages (*datagrams*)  
each datagram's destination independently set

## remote procedure calls

make function calls that run on remote machine

## remote memory access

read/write bytes of data in remote memory

...



# some abstractions

## stream

continuous stream of bytes from one program to another  
'connection' from one program to another

## datagrams

send small messages (*datagrams*)  
each datagram's destination independently set

## remote procedure calls

make function calls that run on remote machine

## remote memory access

read/write bytes of data in remote memory

...

# focus on streams

this class: focus on implementing *streams of bytes*

why?

- most commonly used by applications on the Internet
- many common tasks with other abstractions

# some challenges for streams

separating data into pieces network can handle

putting pieces back together

getting network to send piece to correct remote network

getting network to send piece to correct machine

getting machine to send data to correct program

getting pieces into format wires/radio/fiber/etc. can handle

handling transmission errors

# some challenges for streams

separating data into pieces network can handle

putting pieces back together

getting network to send piece to correct remote network

getting network to send piece to correct remote network  
lots of work! don't want to implement all at once!

getting machine to send data to correct program

getting pieces into format wires/radio/fiber/etc. can handle

handling transmission errors

# some challenges for streams

separating data into pieces network can handle

putting some parts need to be different for different local networks

getting network to send piece to correct remote network

getting network to send piece to correct machine

getting machine to send data to correct program

getting pieces into format wires/radio/fiber/etc. can handle

handling transmission errors

# some challenges for streams

separating some parts should not concern local network implementors

putting pieces back together

getting network to send piece to correct remote network

getting network to send piece to correct machine

getting machine to send data to correct program

getting pieces into format wires/radio/fiber/etc. can handle

handling transmission errors

# some challenges for streams

separating some parts should be same for different abstraction

putting pieces back together

getting network to send piece to correct remote network

getting network to send piece to correct machine

getting machine to send data to correct program

getting pieces into format wires/radio/fiber/etc. can handle

handling transmission errors

# layered model

networking implemented in 'layers'

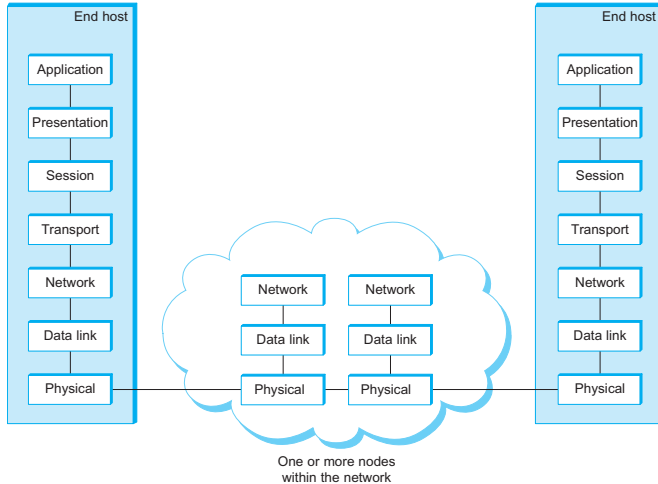
upper layers implemented by making calls to lower layers

example: network implements 'send data to (remote) machine' function ("network layer")

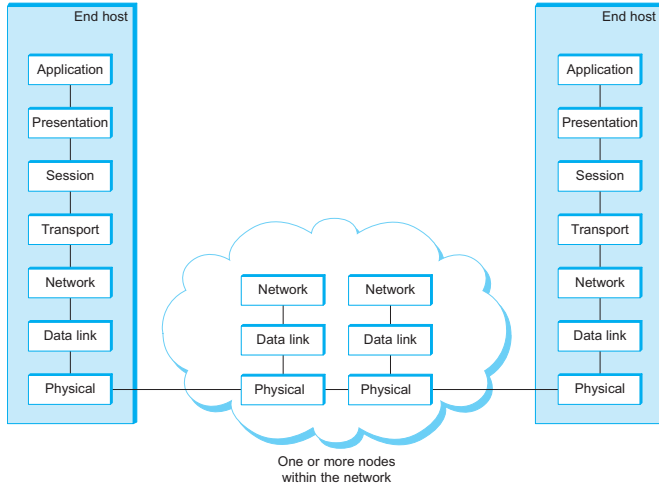
stream implementation calls this to implement 'send stream to remote application'



# OSI model

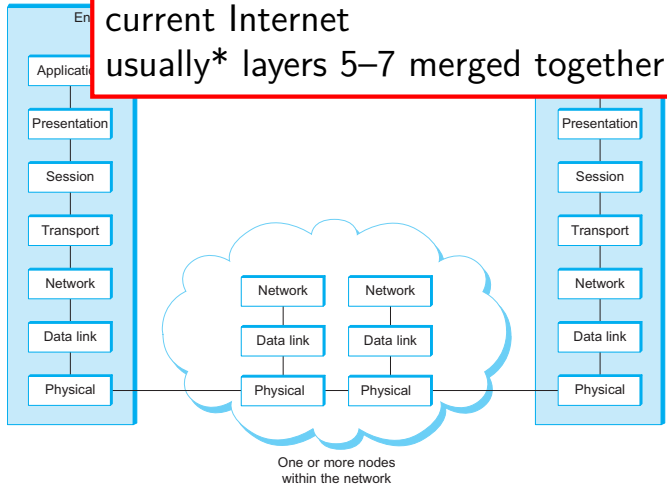


# OSI model



- (7) application:  
what requests/etc.
- (6) presentation:  
data format
- (5) session:  
manage group of streams
- (4) transport:  
streams of data
- (3) network:  
message to correct network
- (2) data link:  
message → bits  
message to correct machine
- (1) physical:  
send bits/...

# OSI model



- (7) **application**:  
what requests/etc.
- (6) **presentation**:  
data format
- (5) **session**:  
manage group of streams
- (4) **transport**:  
streams of data
- (3) **network**:  
message to correct network
- (2) **data link**:  
message → bits  
message to correct machine
- (1) **physical**:  
send bits/...

# OSI model

standardized by ISO (International Standards Organization)

full set of protocols...

- file transfer, message sending, directory lookups ...

that were often implemented and sometimes used...

but mostly lost out to IETF-standardized Internet protocols

- Internet Engineering Task Force

# OSI influence (1)

term 'layer 7', 'layer 4', 'layer 3', etc. almost always refer to OSI model

...even though most of Internet does not follow it  
early Internet protocols predate OSI

## OSI influence (2)

are a lot of Internet protocols influenced by OSI protocols

OSI's DAP (directory access protocol)

adapted into IETF's LDAP (lightweight directory access protocol)

OSI presentation layer ASN.1 used in...

telephony (between telephone companies)

inter-bank messaging

lots of cryptography-related protocols

...

OSI's routing protocol IS-IS still common in large Internet-connected networks

(adapted to work alongside IETF protocols)

# Internet layers

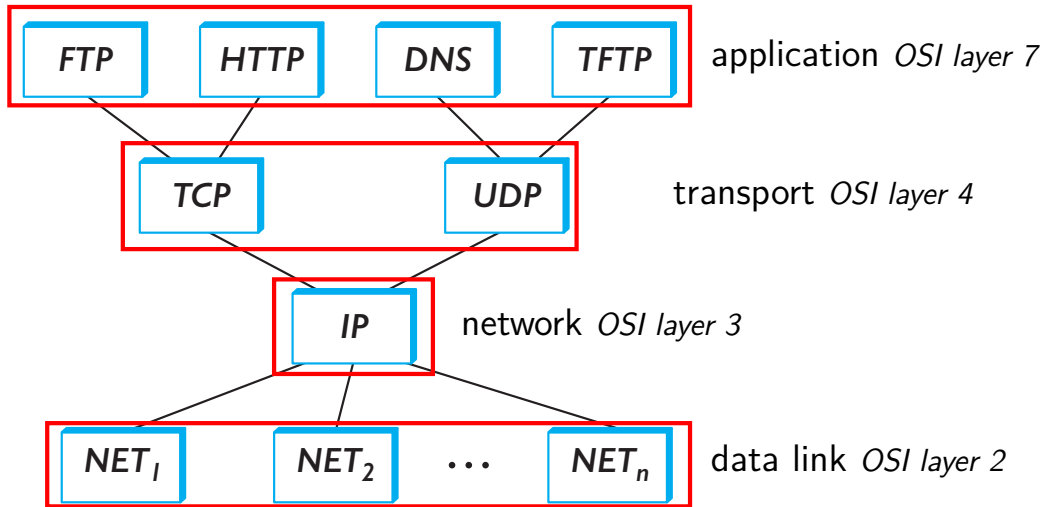
OSI layer	name	examples	purpose
7	application	HTTP, SSH, SMTP, DNS, ...	application-defined meanings
4	transport	TCP, UDP, ...	reach            correct            program, reliability/streams
3	network	IPv4, IPv6, ...	reach            correct            machine (across networks)
2	link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
1	physical	...	encode bits for wire/radio

# Internet layers

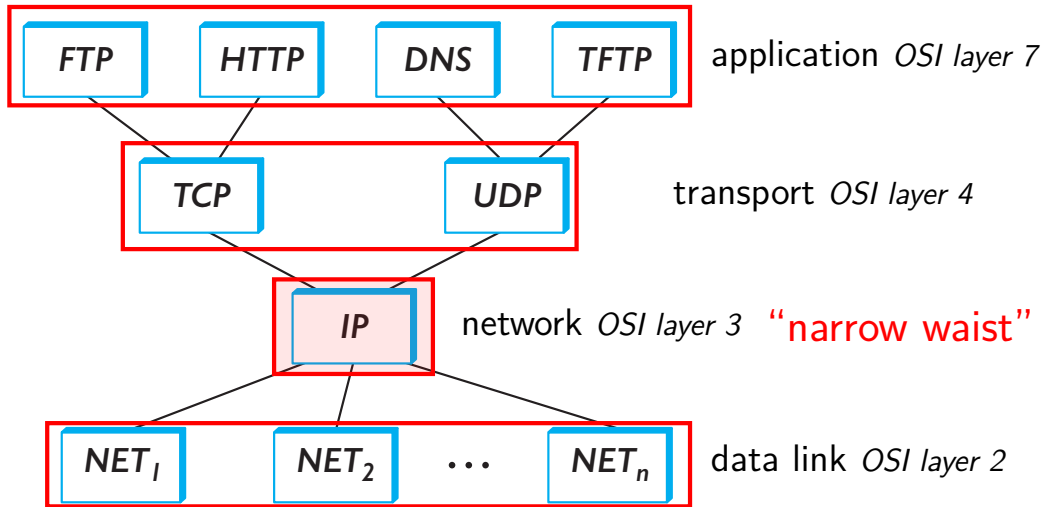
OSI layer	name	examples	purpose
7	application	HTTP, SSH, SMTP, DNS, ...	application-defined meanings
4	transport	TCP, UDP, ...	reach          correct          program, <b>reliability/streams</b>
3	network	IPv4, IPv6, ...	reach          correct          machine (across networks)
2	link	Ethernet, Wi-Fi, ...	coordinate shared wire/radio
1	physical	...	encode bits for wire/radio



# Internet protocols and layers (non-exhaustive)



# Internet protocols and layers (non-exhaustive)



# fuzzy layers (1)

ICMP (Internet Control Message Protocol)...

implemented using a network layer...

so seems like a transport layer protocol?

# fuzzy layers (1)

ICMP (Internet Control Message Protocol)...

implemented using a network layer...

so seems like a transport layer protocol?

used to send errors/control messages about routing...

routing is the network layer's job

so ICMP is part of network layer?

# fuzzy layers (1)

ICMP (Internet Control Message Protocol)...

implemented using a network layer...

so seems like a transport layer protocol?

used to send errors/control messages about routing...

routing is the network layer's job

so ICMP is part of network layer?

I think saying network layer is probably better...

but we're not going to be picky about it

## fuzzy layers (2)

TLS (Transport Control Protocol)...

implemented on top of TCP...

so seems like a application layer protocol?

## fuzzy layers (2)

TLS (Transport Control Protocol)...

implemented on top of TCP...

so seems like a application layer protocol?

used to send other application layer protocols

so maybe a transport layer?

or presentation layer?

I'll call it an application layer...

# ‘extra’ layers

layer terminology doesn't always work cleanly  
often “extra” layers in practice

e.g. HTTPS:

HTTP (app layer) on TLS (another app layer) on TCP (network) on ...

e.g. DNS over HTTPS:

DNS (app layer) on HTTP on on TLS on TCP on ...

e.g. SFTP:

SFTP (app layer??) on SSH (another app layer) on TCP on ...

e.g. HTTP over OpenVPN:

HTTP on TCP on IP on OpenVPN on UDP on different IP on ...



# ‘extra’ layers

layer terminology doesn't always work cleanly  
often “extra” layers in practice

e.g. HTTPS:

HTTP (app layer) on TLS (another app layer) on TCP (network) on ...

e.g. **DNS over HTTPS**:

DNS (app layer) on HTTP on on TLS on TCP on ...

e.g. SFTP:

SFTP (app layer??) on SSH (another app layer) on TCP on ...

e.g. HTTP over OpenVPN:

HTTP on TCP on IP on OpenVPN on UDP on different IP on ...

## protocols usually over HTTP

SOAP (Simple Object Access Protocol) — messaging/remote procedure calls

gRPC (originally from Google) — remote procedure calls

HLS (HTTP Live Streaming) — video streaming

DASH (Dynamic Adaptive Streaming over HTTP) — video streaming

...

# end-to-end argument

Saltzer, Reed, Clark, “End-to-End Arguments in System Design”

“The function in question can completely and correctly be implemented only with the knowledge and help of the application standing at the end points of the communication system.

Therefore, providing that questioned function as a feature of the communication system itself is not possible. (Sometimes an incomplete version of the function provided by the communication system may be useful as a performance enhancement.)”

# end-to-end argument

Saltzer, Reed, Clark, “End-to-End Arguments in System Design”

“The function in question can completely and correctly be implemented **only with the knowledge and help of the application standing at the end points** of the communication system.

Therefore, providing that questioned function as a feature of the communication system itself is not possible. (Sometimes an incomplete version of the function provided by the communication system may be useful as a performance enhancement.)”

## example: reliable file transfer

want to make sure correct data transferred

want to protect against:

- error in hardware/software on sending machine reading file
- bits being flipped in memory on forwarding machine
- communication system flipping bits in data
- hosts crashing during communication

# example: reliable file transfer

want to make sure correct data transferred

want to protect against:

- error in hardware/software on sending machine reading file
- bits being flipped in memory on forwarding machine
- communication system flipping bits in data
- hosts crashing during communication

*communication system can't help a lot of these things*

# example: reliable file transfer

want to make sure correct data transferred

want to protect against:

- error in hardware/software on sending machine reading file

- bits being flipped in memory on forwarding machine

- communication system flipping bits in data

- hosts crashing during communication

*communication system can't help a lot of these things*

*authors experienced router with bad memory/processor*

## **solution: end-to-end checks**

want reliable transfer: compare final files (with hash or similar)

“end-to-end” — doesn't care what middle systems do



# end-to-end argument

Saltzer, Reed, Clark, “End-to-End Arguments in System Design”

“The function in question can completely and correctly be implemented only with the knowledge and help of the application standing at the end points of the communication system.

Therefore, providing that questioned function as a feature of the communication system itself is not possible. (Sometimes an incomplete version of the function provided by the communication system may be useful as a performance enhancement.)”

# end-to-end in practice

“narrow waist” of IP doesn’t provide many guarantees  
no guarantees about reliable transmission, duplicate suppression,  
message order, ...

but try to provide good service (“best effort”)

in design: typically middle systems won’t know/care about what’s  
forwarded

but many exceptions

# backup slides