



# sequence number wraparound

protocol so far requires arbitrarily large sequence numbers


doing  $<$  and  $>$  checks on sequence number, so they need to increase

would like to use smaller sequence numbers


think: transferring multi-gigabyte file

question: what goes wrong when we reuse sequence numbers?


# sender/receiver desync: missing ACKs




# sender/receiver desync: missing ACKs




# sender/receiver desync: missing ACKs




# wraparound



# loss and resend?



# very bad reordering





# possible reason

